

Phase 2 - EDA + Machine Learning with XGB

- More information - Phase 0 (Business Requirement) and Phase 1 (EDA using Python)
- In this Notebook we will do more exploratory data analysis and execute machine learning algorithms
- At the end is going to be provided a summary analysis and expected to answer the business questions (phase 0)

Prediction of Revenue in 2008

Packages used in this Notebook

- tidyverse
- xgboost
- Ckmeans.1d.dp
- caret
- MLmetrics
- to install a package just run the command for example: `install.packages('MLmetrics')`

```
In [2]: library(tidyverse)
library(xgboost)
```

Info - dataset

- the public dataset used in this process can be accessed through IBM website below

<https://www.ibm.com/communities/analytics/watson-analytics-blog/guide-to-sample-datasets/>
(<https://www.ibm.com/communities/analytics/watson-analytics-blog/guide-to-sample-datasets/>)

```
In [3]: df <- read_csv('../data/WA_Retail-SalesMarketing_-ProfitCost.csv')
```

```
## remove records without revenue  
df <- na.omit(df)
```

```
## Columns used for prediction
```

```
col_revenue <- c(  
  'Year'  
  , 'Product line'  
  , 'Product type'  
  , 'Product'  
  , 'Order method type'  
  , 'Retailer country'  
  , 'Revenue')
```

```
df <- df[col_revenue]
```

```
df %>% head(5)
```

Parsed with column specification:

```
cols(  
  Year = col_integer(),  
  `Product line` = col_character(),  
  `Product type` = col_character(),  
  Product = col_character(),  
  `Order method type` = col_character(),  
  `Retailer country` = col_character(),  
  Revenue = col_double(),  
  `Planned revenue` = col_double(),  
  `Product cost` = col_double(),  
  Quantity = col_integer(),  
  `Unit cost` = col_double(),  
  `Unit price` = col_double(),  
  `Gross profit` = col_double(),  
  `Unit sale price` = col_double()  
)
```

Year	Product line	Product type	Product	Order method type	Retailer country	Revenue
2004	Camping Equipment	Cooking Gear	TrailChef Water Bag	Telephone	United States	315044.33
2004	Camping Equipment	Cooking Gear	TrailChef Water Bag	Telephone	Canada	13444.68
2004	Camping Equipment	Cooking Gear	TrailChef Water Bag	Telephone	Japan	181120.24
2004	Camping Equipment	Cooking Gear	TrailChef Water Bag	Telephone	China	69608.15
2004	Camping Equipment	Cooking Gear	TrailChef Water Bag	Telephone	Singapore	30940.35

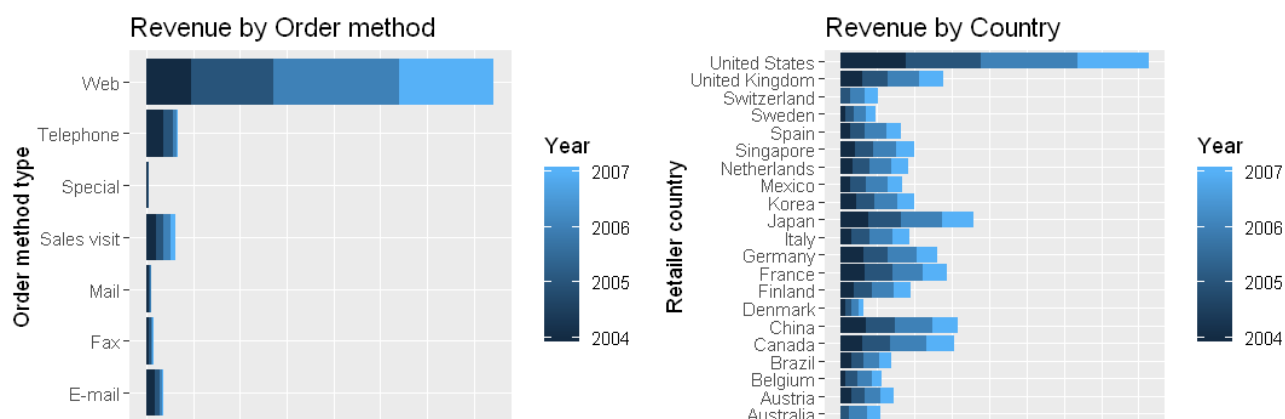
Basic information / distribution of the data

```
In [4]: df <- df %>% mutate_if(is.character, as.factor)

options(repr.plot.width = 9, repr.plot.height = 3)
gg1 <- qplot(data=df, x=`Order method type`, y=Revenue, geom='col', fill=Year, main =
'Revenue by Order method' ) +
  scale_y_continuous(name="Revenue", labels = scales::comma) + coord_flip() +
  theme(axis.title.x=element_blank(),
        axis.text.x=element_blank(),
        axis.ticks.x=element_blank())

gg2 <- qplot(data=df, x=`Retailer country`, y=Revenue, geom='col', fill=Year, main =
'Revenue by Country' ) +
  scale_y_continuous(name="Revenue", labels = scales::comma) + coord_flip() +
  theme(axis.title.x=element_blank(),
        axis.text.x=element_blank(),
        axis.ticks.x=element_blank())

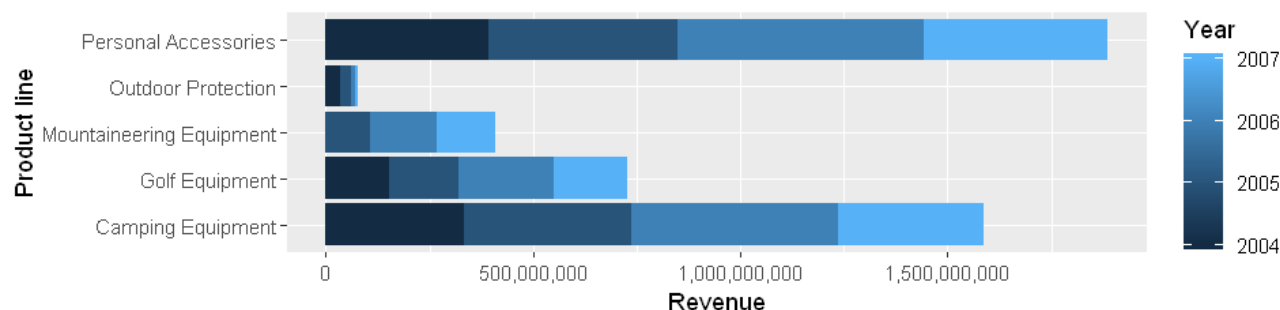
gridExtra::grid.arrange(gg1, gg2, nrow = 1)
```



Analyze Revenue by Product line with Year distribution

```
In [5]: options(repr.plot.width = 8, repr.plot.height = 2)

## revenue by product line by year
qplot(data=df, x=`Product line`, y=Revenue, geom='col', fill=Year) +
  scale_y_continuous(name="Revenue", labels = scales::comma) + coord_flip()
```

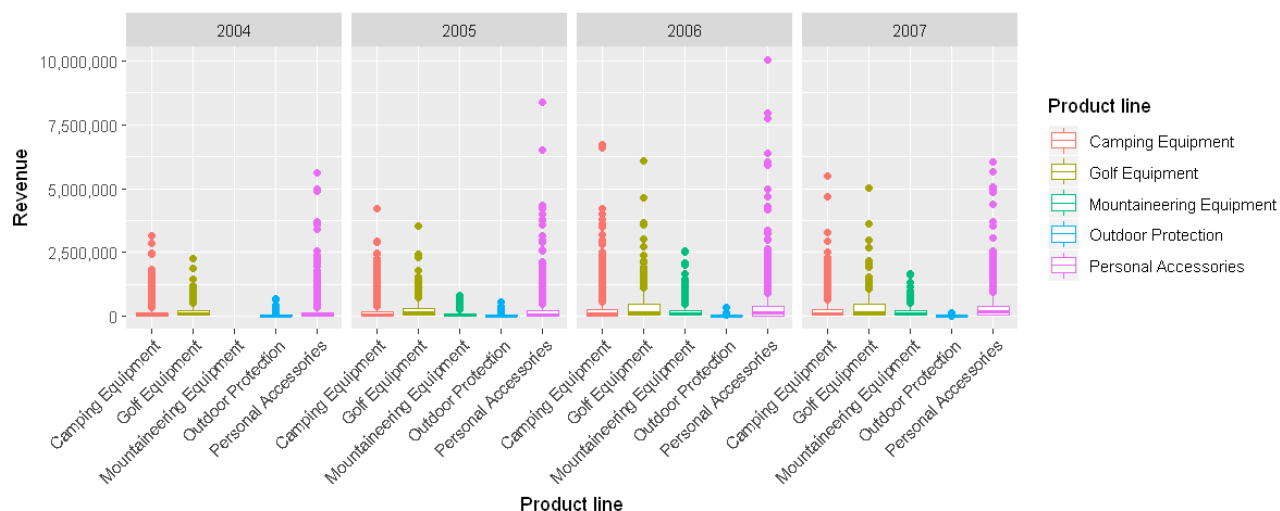


1st Note

- Personal Accessories and Camping Equipment tend to have higher revenue than other Product lines
- Web method provides the biggest revenues
- United States country also provides the biggest revenue

Distribution of the Revenue and evaluation of possible outliers

```
In [6]: options(repr.plot.width = 10, repr.plot.height = 4)
qplot(data=df, x=`Product line`, y=Revenue, geom='boxplot', color=`Product line`, fac
ets = ' Year ~ .') +
  scale_y_continuous(name="Revenue", labels = scales::comma) + ## mediana = 163000
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



2nd Note

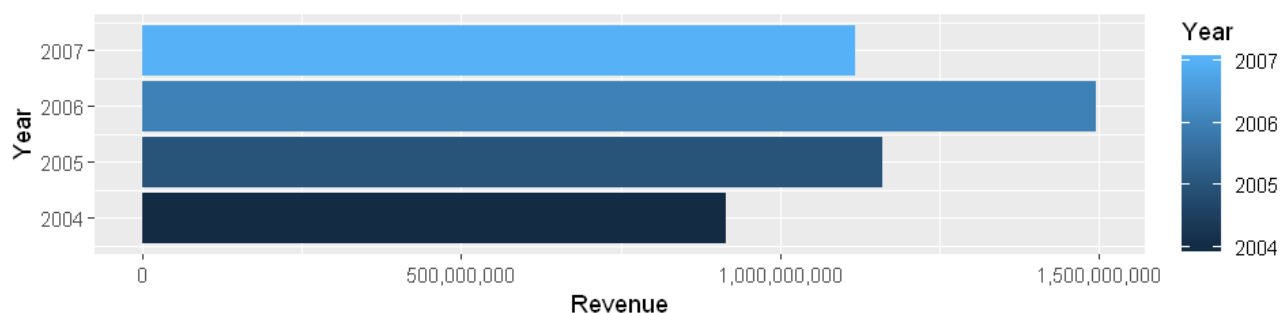
- Personal Accessories and Camping Equipment have more outliers with higher value in Year 2006
- Let's confirm the Total Revenue by Year (2006) and it's impact related to others

Info

- this type of scenario (higher revenue for specific items on specific year) must be better understanding with Line of Business Managers and also evaluate if it was just a sazonal discrepance or some marketing campaign or business decision have bigger influence on these numbers
- Another important point to understand are related on how it can influence the prediction of the revenue for year 2008

Total Revenue by Year

```
In [7]: ## revenue by year -> fazer a simulacao/modelo somente com estes dados
options(repr.plot.width = 8, repr.plot.height = 2)
qplot(data=df, x=Year, y=Revenue, geom='col', fill=Year)+
  scale_y_continuous(name="Revenue", labels = scales::comma) + coord_flip()
```



Start the Prediction of Revenue for year - 2008

- Will be used XGBoost to predict the revenue in 2008
- Use Caret to partition the train and test data
- Compare the results achieved with XGB vs Random Forest in Phase 1

Additional info

- Were not provided one dataset by the business users related to year 2008

XGBoost documentation and samples:

<https://xgboost.readthedocs.io/en/latest/> (<https://xgboost.readthedocs.io/en/latest/>)

<https://github.com/dmlc/xgboost> (<https://github.com/dmlc/xgboost>)

```
In [8]: ## define the target and prepare the data to run XGB
target <- 'Revenue'
idx_target <- which(col_revenue == target)

dfxgb <- df %>% mutate_if(is.character, as.factor)
dfxgb <- dfxgb %>% mutate_if(is.factor, as.integer)

set.seed(12345)
idx <- caret::createDataPartition(dfxgb$Revenue, p=0.80, list=FALSE)
train <- dfxgb[idx, ]
test <- dfxgb[-idx, ]

dtrain <- xgb.DMatrix(data = as.matrix(train[, -idx_target]), label= train[[target]])
dtest <- xgb.DMatrix(data = as.matrix(test[, -idx_target]))
```

```
In [9]: ## execution of XGB
set.seed(12345)
fit.xgb <- xgboost(data = dtrain,
                   objective = "reg:linear",
                   booster = "gbtree",
                   print_every_n = 25, ## print every 25
                   nrounds = 350)
```

```
[1]      train-rmse:368519.000000
[26]      train-rmse:161832.156250
[51]      train-rmse:131686.859375
[76]      train-rmse:110794.734375
[101]     train-rmse:101309.117188
[126]     train-rmse:92665.507812
[151]     train-rmse:87156.484375
[176]     train-rmse:83507.718750
[201]     train-rmse:79411.625000
[226]     train-rmse:75475.070312
[251]     train-rmse:72690.914062
[276]     train-rmse:70538.992188
[301]     train-rmse:68313.281250
[326]     train-rmse:66531.570312
[350]     train-rmse:64893.933594
```

3rd Note

- XGB have a lot of parameters to tune the algorithm as can be seen in the XGB documentation link provided above
- The printed result above shows the improvement of the training process at every 25 steps as detailed in `print_every_n = 25`, and use `rmse` evaluation metric

Info

- The Random Forest Regressor/Phase_1 (Python - sklearn) use R^2 to evaluate the prediction and will be used also below to compare both metrics. Our target is to achieve R^2 near to 1

RMSE : Root Mean Squared Error

https://en.wikipedia.org/wiki/Root-mean-square_deviation (https://en.wikipedia.org/wiki/Root-mean-square_deviation)

R^2 : Coefficient of determination

https://en.wikipedia.org/wiki/Coefficient_of_determination (https://en.wikipedia.org/wiki/Coefficient_of_determination)

** In regression, the R^2 coefficient of determination is a statistical measure of how well the regression predictions approximate the real data points. An R^2 of 1 indicates that the regression predictions perfectly fit the data

Feature importance

Note

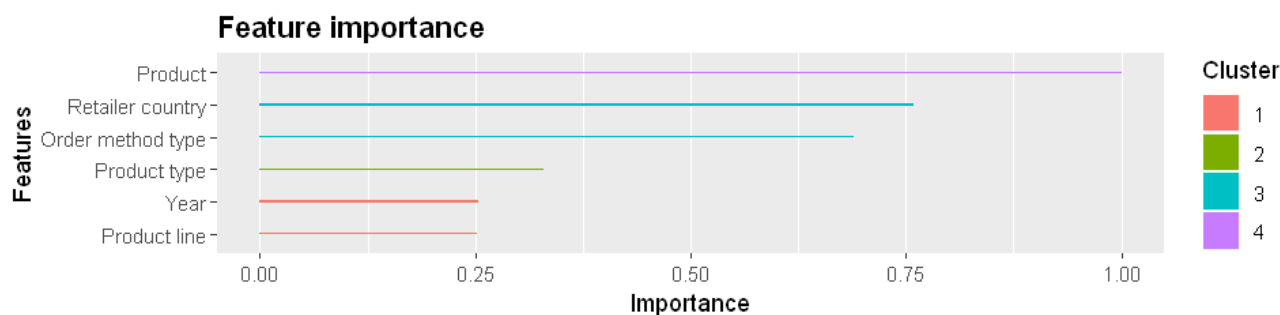
- the 3 most important features still the same (Product, Retailer country and Order method type) however with XGB the most important feature is Product

```
In [10]: ### feature importance
## importance variable of the xgb model
imp_features <- xgb.importance(model = fit.xgb)

print(imp_features)

## Plot Feature importance
xgb.ggplot.importance(imp_features, rel_to_first = TRUE)
```

	Feature	Gain	Cover	Frequency
1:	Product	0.30438392	0.49273598	0.25324370
2:	Retailer country	0.23113830	0.22153642	0.30001880
3:	Order method type	0.21003098	0.08815150	0.17826250
4:	Product type	0.10035484	0.09174638	0.06788266
5:	Year	0.07737424	0.07443256	0.15739000
6:	Product line	0.07671772	0.03139716	0.04320233



Compare the prediction between XGB and Random Forest (Phase_1)

```
In [11]: ## Evaluate the prediction with test data
predict_xgb <- predict(fit.xgb, dtest)
predict_xgb <- ifelse(predict_xgb < 1, 0, predict_xgb)

## R^2 0.9016625 -> 90% -> 6% better than random forest python (84%)
print('----- R^2 evaluation')
print(MLmetrics::R2_Score(predict_xgb, test$Revenue))

[1] "----- R^2 evaluation"
[1] 0.9016625
```

Important note

- The prediction provided by XGB achieve an R^2 of 0.9016625 and is 6% higher than Random Forest in Phase 1 (0.8442048767542322) and is a good improvement, so we can keep the prediction with XGB
- Let's move on and predict and evaluate the revenue for the year 2008

```
In [12]: ## The Business Requirement request to predict the revenue for 2008 but do not provide data for 2008

## the trick here => we will use the same products sold on 2007 to predict the Revenue for 2008 and evaluate the results
#### the model used have R^2 of 0.9016 so we expect to achieve a confident result

df_2008 <- dfxgb %>% filter(Year==2007)

## update Year to 2008 and setup the Revenue to 0 -> Revenue will be updated with the prediction later
df_2008$Year <- 2008
df_2008$Revenue <- 0

print('----- INITIAL dataset: 2008 YEAR')
print(summary(df_2008))

## generate the data to use with XGB model (fit.xgb)
xgb_2008 <- xgb.DMatrix(data = as.matrix(df_2008[, -idx_target]), label= df_2008[[target]])

predict_revenue_2018 <- predict(fit.xgb, xgb_2008)
predict_revenue_2018 <- ifelse(predict_revenue_2018 < 1, 0, predict_revenue_2018)

## update the Revenue prediction for year 2008 and compare with other Years
df_2008$Revenue <- predict_revenue_2018

## generate one dataset to compare the revenue form all Years
df_all <- rbind(dfxgb, df_2008)

print('----- PREDICTION: 2008 YEAR')
summary(df_all)
```

```
[1] "----- INITIAL dataset: 2008 YEAR"
```

Year	Product line	Product type	Product
Min. :2008	Min. :1.000	Min. : 1.00	Min. : 1.00
1st Qu.:2008	1st Qu.:1.000	1st Qu.: 4.00	1st Qu.: 36.00
Median :2008	Median :3.000	Median :10.00	Median : 68.00
Mean :2008	Mean :2.937	Mean :10.64	Mean : 70.36
3rd Qu.:2008	3rd Qu.:5.000	3rd Qu.:16.00	3rd Qu.:106.00
Max. :2008	Max. :5.000	Max. :21.00	Max. :144.00
Order method type	Retailer country	Revenue	
Min. :1.000	Min. : 1.00	Min. :0	
1st Qu.:4.000	1st Qu.: 6.00	1st Qu.:0	
Median :7.000	Median :12.00	Median :0	
Mean :5.773	Mean :11.57	Mean :0	
3rd Qu.:7.000	3rd Qu.:17.00	3rd Qu.:0	
Max. :7.000	Max. :21.00	Max. :0	

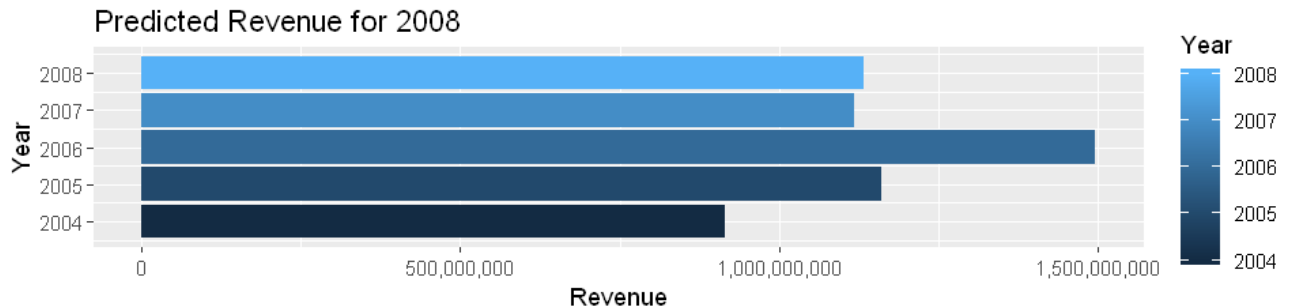
```
[1] "----- PREDICTION: 2008 YEAR"
```

Year	Product line	Product type	Product
Min. :2004	Min. :1.000	Min. : 1.00	Min. : 1.00
1st Qu.:2005	1st Qu.:1.000	1st Qu.: 4.00	1st Qu.: 35.00
Median :2006	Median :3.000	Median :10.00	Median : 71.00
Mean :2006	Mean :2.951	Mean :10.52	Mean : 71.63
3rd Qu.:2007	3rd Qu.:5.000	3rd Qu.:16.00	3rd Qu.:109.00
Max. :2008	Max. :5.000	Max. :21.00	Max. :144.00
Order method type	Retailer country	Revenue	
Min. :1.000	Min. : 1.00	Min. : 0	
1st Qu.:4.000	1st Qu.: 6.00	1st Qu.: 19549	
Median :6.000	Median :11.00	Median : 65002	
Mean :5.309	Mean :11.46	Mean : 197870	
3rd Qu.:7.000	3rd Qu.:17.00	3rd Qu.: 207914	
Max. :7.000	Max. :21.00	Max. :10054289	

Evaluate the Revenue Prediction for 2008

- Info: The Revenue for 2008 Year are predictions and from 2004 to 2007 are real Revenue

```
In [13]: ## info: 2008 is the Revenue Prediction
qplot(data=df_all, x=Year, y=Revenue, geom='col', fill=Year, main='Predicted Revenue
for 2008')+
  scale_y_continuous(name="Revenue", labels = scales::comma) + coord_flip()
```



```
In [14]: ## Revenue by Year
revenue_by_year <- df_all %>% group_by(Year) %>% summarise(Total_Revenue = sum(Revenue))
print(revenue_by_year)
```

```
# A tibble: 5 x 2
  Year Total_Revenue
<dbl>         <dbl>
1  2004    914352804.
2  2005   1159195590.
3  2006   1495891101.
4  2007   1117336274.
5  2008   1132382363.
```

Insights

- The revenue prediction for the year 2008 is bigger than 2007 and smaller than the year 2005
- The higher revenue of year 2006 seems to not impact the prediction of the revenue for 2008
 - => the higher outliers exposed in the boxplot chart seems to not impact in the revenue prediction
- The confidence of this prediction is quite high so let's answer the final business questions
 - what are the divergence expected on revenue in 2008 ?

Analysis - Revenue Prediction of 2008 vs (2005 and 2007)

```
In [15]: rev_2008_vs_2005 <- round( (revenue_by_year[revenue_by_year$Year==2008, 'Total_Revenue'] /
                                     revenue_by_year[revenue_by_year$Year==2005, 'Total_Revenue']) * 100 , 2)

rev_2008_vs_2007 <- round( (revenue_by_year[revenue_by_year$Year==2008, 'Total_Revenue'] /
                             revenue_by_year[revenue_by_year$Year==2007, 'Total_Revenue']) * 100 , 2)

print(' ----- Revenue 2008 vs 2005 ')
colnames(rev_2008_vs_2005) <- 'Percent Revenue_2008_vs_2005'
rev_2008_vs_2005

print(' ----- Revenue 2008 vs 2007 ')
colnames(rev_2008_vs_2007) <- 'Percent Revenue_2008_vs_2007'
rev_2008_vs_2007
```

```
[1] " ----- Revenue 2008 vs 2005 "
```

Percent Revenue_2008_vs_2005
97.69

```
[1] " ----- Revenue 2008 vs 2007 "
```

Percent Revenue_2008_vs_2007
101.35

Summary

- Answers for the Business Requirement

The revenue prediction for year 2008 can be done and the 3 most important features to predict revenue are

- Product
- Retailer country and
- Order method type

The prediction confidence is high

- Achieved an R^2 of 0.9016 and we can fill confident with the prediction numbers

The divergence related to the revenue prediction are

** Based on the assumption to sell the same products of the year 2007

- Revenue prediction for 2008 are expected to be 2.5% smaller than the year 2005 and
- Revenue prediction for 2008 are expected to be 1.5% higher than the last year (year 2007)

```
In [16]: print('----- THE END')
```

```
[1] "----- THE END"
```