

# Phase 2 - Evaluate the Recall metric and analyze the unbalanced class

## Summary of all steps executed in this notebook

- Start loading the 3 exported models (xgb and gbm) with higher accuracy and run the predictions. These models were exported in previous notebooks and were built using the ML frameworks sklearn and h2o.ai
- Build new machine learning models with the additional techniques explained below and
- Evaluate the machine learning model with higher RECALL metric for deployment  
Recall metric is one good option to evaluate this type of scenario (customer churn)  
Recall is also known as Sensitivity or True positive rate(TPR)

## Additional techniques trying to improve the performance of the ML model, such as:

- Evaluation of unbalanced classification / class weight
- Smote technique for oversampling the training dataset
- Apply Standard Scale in the data before build the ML models, and compare with previous results using default data format

## Quick recap of previous notebooks - Machine Learning with Python ¶

### Build 15 machine learning using different frameworks - XGBoost, LightGBM, Sklearn, H2O.ai and Apache Spark

- 2 ML models in 1x5
- 6 ML models in 2x6
- 4 ML models in 3x6
- 3 ML models in 4x6

### Model evaluation

- Previous notebooks use the global metric accuracy for evaluation

### Starting process... load python modules and also connect to h2o cluster

In [1]:

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

## Remove warnings
import warnings
warnings.filterwarnings("ignore")

## Machine Learning Models
import h2o
from h2o.estimators import H2OXGBoostEstimator, H2OGradientBoostingEstimator
import xgboost as xgb
from sklearn.ensemble import GradientBoostingClassifier

## Sklearn Metrics - Classification and model selection
from sklearn.metrics import (confusion_matrix, classification_report, accuracy_score,
                             roc_auc_score, recall_score, roc_auc_score)

from sklearn.model_selection import train_test_split

## Scale data
from sklearn.preprocessing import scale

# import imblearn - SMOTE
from imblearn.over_sampling import SMOTE

## Load ML models
import pickle

## connect to h2o cluster and remove all object
h2o.connect(ip='192.168.56.102')
h2o.remove_all()
```

Connecting to H2O server at http://192.168.56.102:54321 ... successful.

|                               |   |
|-------------------------------|---|
| H2O cluster uptime:           | 5 hours 38 mins   |
| H2O cluster timezone:         | America/Sao_Paulo   |
| H2O data parsing<br>timezone: | UTC   |
| H2O cluster version:          | 3.26.0.3  |
| H2O cluster version<br>age:   | 7 months and 23 days !!!                                  |
| H2O cluster name:             | users1  |
| H2O cluster total nodes:      | 1   |
| H2O cluster free<br>memory:   | 5.746 Gb  |
| H2O cluster total cores:      | 1   |
| H2O cluster allowed<br>cores: | 1   |
| H2O cluster status:           | locked, healthy   |
| H2O connection url:           | http://192.168.56.102:54321                               |
| H2O connection proxy:         | None  |
| H2O internal security:        | False   |
| H2O API Extensions:           | Amazon S3, XGBoost,<br>Algos, AutoML, Core<br>V3, Core V4 |
| Python version:               | 3.7.3 final   |

## Load and prepare the dataset to build ML models

- H2O - pandas Dataframe with raw data
- Sklearn - Numpy arrays with One Hot Encode for categorical columns

In [2]:

```

## Load and prepare the dataset to build ML models - Customer churn
df = pd.read_csv('../data/WA_Fn-UseC_-Telco-Customer-Churn.csv')
df.loc[(df.tenure==0) & (df.TotalCharges == ' '), ['TotalCharges', 'tenure']] = 0
df['TotalCharges'] = df['TotalCharges'].astype('float')

current_features = ['tenure', 'MonthlyCharges', 'TotalCharges', 'gender', 'PaymentMethod', 'Churn', 'Contract']

df = df[current_features]

## H2O DATA - FORMAT AS DATA FRAME WITHOUT ONE HOT ENCODE
target = 'Churn'
features_h2o = df.columns.to_list()
features = features_h2o.copy()
features.remove(target)

X = df[features].copy()
y = df[target]

SEED = 42
X_train_h2o, X_test_h2o, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=SEED)

X_train_h2o[target] = y_train
X_test_h2o[target] = y_test

## Convert to h2o Frame
train_h2o = h2o.H2OFrame(X_train_h2o, destination_frame='train.hex')
train_h2o[target] = train_h2o[target].asfactor()
test_h2o = h2o.H2OFrame(X_test_h2o, destination_frame='test.hex')
test_h2o[target] = test_h2o[target].asfactor()

## SKLEARN with One Hot Encode and format as Numpy Array
## One Hot Encode for categorical features - sklearn

target_1_0 = lambda x: ['No', 'Yes'].index(x)
df[target] = df[target].apply(target_1_0)

OHE_cols = ['gender', 'PaymentMethod', 'Contract']
df = pd.get_dummies(data=df, columns=OHE_cols)

features = df.columns.to_list()
features.remove(target)

X = df[features].values
y = df[target].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=SEED)

df.head(2)

```

Out[2]:

◀ [REDACTED] ▶

In [3]:

## Export samples of data for new predictions

- In [4]:

## Machine Learning models

- file:///F:/ D HDD/R Python Spark/GITHUB PORTFOLIO IN/ GIT CUSTOMER SUPPORT/DataScience projects-master/DataScience proj... 5/19

In [5]:

```

## Function to print Confusion Matrix and metrics
def rpt_metrics_report(y_true, y_pred, msg_model='MODEL XPTO', rpt_confusion_matrix=False):
    """Print metrics """

    accuracy_score_rpt = accuracy_score(y_true, y_pred)
    recall_score_rpt = recall_score(y_true, y_pred)
    auc_rpt = roc_auc_score(y_true, y_pred)

    print('Model: ', msg_model)
    print('-- Accuracy: ', accuracy_score_rpt)
    print('-- AUC      : ', auc_rpt)
    print('-- Recall  : ', recall_score_rpt)
    print('')
    if rpt_confusion_matrix:
        report = classification_report(y_true, y_pred)
        confusion_matrix_rpt = confusion_matrix(y_true, y_pred)
        print('-- Confusion Matrix')
        print('0  FP')
        print('FN 1')
        print('')
        print(confusion_matrix_rpt)
        print('')
        print('')
        print('-- Metrics report')
        print(report)
        print('')

```

## Load the machine learning models already build and start the evaluation using the metric recall

### GBM and XBG (sklearn and h2o)

- These 3 models presented below have similar accuracy score and were the ones that had the best results  
exported previously in specific notebook

In [6]:

```

## Load models and validate
print('')
print('ML prediction - sklearn baseline models')
print('')

## GBM model
file_export_model = './ML_models/model_GBM_baseline_v1.sav'
gbm_model = pickle.load(open(file_export_model, 'rb'))
y_pred_gbm = gbm_model.predict(X_test)
rpt_metrics_report(y_test, y_pred_gbm, 'GBM')

## XGB model
file_export_model = './ML_models/model_XGB_baseline_v1.sav'
xgb_model = pickle.load(open(file_export_model, 'rb'))
y_pred_xgb = xgb_model.predict(X_test)
rpt_metrics_report(y_test, y_pred_xgb, 'XGB')

## H2O baseline models
## get model exported - predict and generate target in numeric format

## XGB model
export_model_path = './ML_models/model_xgb_v1/fit_xgb.model"
h2o_model_xgb = h2o.load_model(export_model_path)
# model_xgb.model_performance()
ypred_XGB_df = h2o_model_xgb.predict(test_h2o).as_data_frame()
y_pred_h2o = ypred_XGB_df['predict'].apply(target_1_0)
y_true_test_h2o = X_test_h2o[target].apply(target_1_0).values
rpt_metrics_report(y_true_test_h2o, y_pred_h2o, 'H2O - XGB')

```

ML prediction - sklearn baseline models

Model: GBM

```

-- Accuracy: 0.8
-- AUC      : 0.7095220490261945
-- Recall   : 0.5127388535031847

```

Model: XGB

```

-- Accuracy: 0.7991397849462366
-- AUC      : 0.7019112033599947
-- Recall   : 0.49044585987261147

```

xgboost prediction progress: |

| 100%

Model: H2O - XGB

```

-- Accuracy: 0.7918279569892474
-- AUC      : 0.7580903355115246
-- Recall   : 0.6847133757961783

```

## 1st Machine learning model insight

- The 3 models have similar accuracy between 79% and 80%, however the model with H2O (xgb) has an AUC 5% higher
  - Recall with xgb using H2O.ai framework is 68,47% => 17% higher than the others
- The H2O XGB is currently the model that provide the best performance in general

The difference between these 2 executions are mainly related to default hyper parameter used by each framework

(ML model and framework - sklearn vs. h2o.ai)

Some examples:

- GBM with sklearn use learning rate = 0.1 and xgb with H2O.ai use eta/learning rate = 0.3 as can be seen below.
- Number of trees, and so on...

### GBM with sklearn

In [7]:

```
gbm_model
```

Out[7]:

```
GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=N
one,
                           learning_rate=0.1, loss='deviance', max_depth=
3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=N
one,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           n_iter_no_change=None, presort='deprecated',
                           random_state=42, subsample=1.0, tol=0.0001,
                           validation_fraction=0.1, verbose=0,
                           warm_start=False)
```

### XGB with H2O.ai



In [8]:

```
h2o_model_xgb.default_params
```

Out[8]:

```
{'model_id': None,
 'training_frame': None,
 'validation_frame': None,
 'nfolds': 0,
 'keep_cross_validation_models': True,
 'keep_cross_validation_predictions': False,
 'keep_cross_validation_fold_assignment': False,
 'score_each_iteration': False,
 'fold_assignment': 'AUTO',
 'fold_column': None,
 'response_column': None,
 'ignored_columns': None,
 'ignore_const_cols': True,
 'offset_column': None,
 'weights_column': None,
 'stopping_rounds': 0,
 'stopping_metric': 'AUTO',
 'stopping_tolerance': 0.001,
 'max_runtime_secs': 0.0,
 'seed': -1,
 'distribution': 'AUTO',
 'tweedie_power': 1.5,
 'categorical_encoding': 'AUTO',
 'quiet_mode': True,
 'export_checkpoints_dir': None,
 'ntrees': 50,
 'max_depth': 6,
 'min_rows': 1.0,
 'min_child_weight': 1.0,
 'learn_rate': 0.3,
 'eta': 0.3,
 'sample_rate': 1.0,
 'subsample': 1.0,
 'col_sample_rate': 1.0,
 'colsample_bylevel': 1.0,
 'col_sample_rate_per_tree': 1.0,
 'colsample_bytree': 1.0,
 'max_abs_leafnode_pred': 0.0,
 'max_delta_step': 0.0,
 'monotone_constraints': None,
 'score_tree_interval': 0,
 'min_split_improvement': 0.0,
 'gamma': 0.0,
 'nthread': -1,
 'max_bins': 256,
 'max_leaves': 0,
 'min_sum_hessian_in_leaf': 100.0,
 'min_data_in_leaf': 0.0,
 'sample_type': 'uniform',
 'normalize_type': 'tree',
 'rate_drop': 0.0,
 'one_drop': False,
 'skip_drop': 0.0,
 'tree_method': 'auto',
 'grow_policy': 'depthwise',
 'booster': 'gbtree',
 'reg_lambda': 1.0,
 'reg_alpha': 0.0,
 'dmatrix_type': 'auto',
```

```
'backend': 'auto',
'gpu_id': 0}
```

**To finalize the evaluation and decide which model to deploy in production additional techniques will be applied**

- Apply Standard Scale in the data before build the machine learning model
- Evaluate the balance of target class and optimize the model execution

## Evaluation of target class distribution

- Churn : ~ 26%%
- No churn : ~ 73%

In [9]:

```
print('Target class distribution - all data')
print(pd.value_counts(df[target]))
print(pd.value_counts(df[target])/len(df[target]))
print('')
print('Target class distribution - train dataset')
print(pd.value_counts(y_train))
print(pd.value_counts(y_train)/len(y_train))
```

Target class distribution - all data

0 5174

1 1869

Name: Churn, dtype: int64

0 0.73463

1 0.26537

Name: Churn, dtype: float64

Target class distribution - train dataset

0 3477

1 1241

dtype: int64

0 0.736965

1 0.263035

dtype: float64

## Machine learning models with scale and class weight

### XGB - sklearn

- Standard Scale applied and scale\_pos\_weight = class weight
- Recall achieved : 82,0%

In [10]:

```
class_weight = sum(y_train == 0) / sum(y_train == 1) ## total 0:3477 / total 1: 1241 =
2.8017727639000807
print('Class weight: ', class_weight)
```

Class weight: 2.8017727639000807

In [11]:

```
## XGB with class weight - scale_pos_weight

## Scale data - only tenure and charges (numeric values)
scale_trainX = scale(X_train[:, [0, 1, 2]])
scale_trainX = np.concatenate((scale_trainX, X_train[:, 3:12 ]), axis = 1)

scale_testX = scale(X_test[:, [0, 1, 2]])
scale_testX = np.concatenate((scale_testX, X_test[:, 3:12 ]), axis = 1)

model_xgb_scale = xgb.XGBClassifier(random_state=SEED, scale_pos_weight=class_weight)
model_xgb_scale.fit(scale_trainX, y_train)

y_pred_xgb_scale = model_xgb_scale.predict(scale_testX)
rpt_metrics_report(y_test, y_pred_xgb_scale, 'XGB - CLASS WEIGHT')
```

Model: XGB - CLASS WEIGHT

```
-- Accuracy: 0.7436559139784946
-- AUC      : 0.7677218883830214
-- Recall   : 0.8200636942675159
```

## 2nd Insight

- The accuracy score is 5% worse however the recall metric is much higher
- XGB with scale and class weight achieved 82% of recall

## GBM execution with Standard Scale and SMOTE to balance the class - Churn

- GBM with sklearn does not have the option to inform the class weight to the algorithm, so let's balance the class with SMOTE

### GBM - model evaluation

- Recall of 86,30%

In [12]:

```
# SMOTE with imblearn
from imblearn.over_sampling import SMOTE
smote_bal = SMOTE(random_state=SEED)
X_smote, y_smote = smote_bal.fit_resample(scale_trainX, y_train)

## GBM smote
model_GBM_smote = GradientBoostingClassifier(random_state=SEED)
model_GBM_smote.fit(X_smote, y_smote)

y_pred_gbm_smote = model_GBM_smote.predict(scale_testX)
rpt_metrics_report(y_test, y_pred_gbm_smote, 'GBM - scale and SMOTE ')

print('Balance class - Churn')
print(pd.value_counts(y_smote))
```

```
Model: GBM - scale and SMOTE
-- Accuracy: 0.7290322580645161
-- AUC      : 0.7712458103284553
-- Recall   : 0.8630573248407644
```

```
Balance class - Churn
1    3477
0    3477
dtype: int64
```

## Execute the same process with XGB without class weight but with SMOTE and standard scale

- Recall: 85,50% and a bit small than GBM

In [13]:

```
## GBM smote
model_XGB_smote = xgb.XGBClassifier(random_state=SEED)
model_XGB_smote.fit(X_smote, y_smote)

y_pred_xgb_smote = model_XGB_smote.predict(scale_testX)
rpt_metrics_report(y_test, y_pred_xgb_smote, 'XGB - scale and SMOTE ')
```

```
Model: XGB - scale and SMOTE
-- Accuracy: 0.7286021505376344
-- AUC      : 0.7684434689917389
-- Recall   : 0.8550955414012739
```

## GBM - model evaluation - sample with hyper parameter optimization

- Recall: 88,05% => almost 2% higher compared to default GBM parameters
- Export gbm model for deployment - best recall metric

```
## tune GBM sample
model_GBM_smote = GradientBoostingClassifier(random_state=SEED, learning_rate=0.3, n_estimators=200, max_depth=5)
model_GBM_smote.fit(X_smote, y_smote)

y_pred_gbm_smote = model_GBM_smote.predict(scale_testX)
rpt_metrics_report(y_test, y_pred_gbm_smote, 'GBM - scale and SMOTE ')

## Export model for production deployment
file_export_model = './ML_models/model_GBM_prod_v1.sav'
pickle.dump(model_GBM_smote, open(file_export_model, 'wb'))
print("Model exported")
```

Model exported

- Xgb in h2o framework does not have in version 3.26 the option of scale weight or balance\_class however other ML algorithms have the option to inform balance\_classes=True
- Execution of GBM and GLM (2 models as example) just to evaluate the results with balance class = True

```
from h2o.estimators import H2OGradientBoostingEstimator, H2OGeneralizedLinearEstimator,
H2ODeepLearningEstimator
```

```
## H2O SMOTE AND BALANCE
hfeatures = features_h2o.copy()
hfeatures.remove(target)
glm_h2o_smote_cw = H2OGeneralizedLinearEstimator(balance_classes=True, seed=SEED, family='binomial')
glm_h2o_smote_cw.train(
    x = hfeatures,
    y = target,
    training_frame = train_h2o)
y_pred_glm_smote = glm_h2o_smote_cw.predict(test_h2o).as_data_frame()
y_pred_glm_h2o = y_pred_glm_smote['predict'].apply(target_1_0)
rpt_metrics_report(y_true_test_h2o, y_pred_glm_h2o, 'H2O - GLM - SCALE AND BALANCE CLASS')

```

file:///F:/ D HDD/R Python Spark/GITHUB PORTFOLIO IN/ GIT CUSTOMER SUPPORT/DataScience projects-master/DataScience pr... 14/19



In [18]:

```
## GBM model

## Baseline model v1
file_export_model = './ML_models/model_GBM_baseline_v1.sav'
gbm_model = pickle.load(open(file_export_model, 'rb'))
y_pred_gbm = gbm_model.predict(X_test)
rpt_metrics_report(y_test, y_pred_gbm, 'GBM', rpt_confusion_matrix=True)

## GBM Model for production
file_export_model = './ML_models/model_GBM_prod_v1.sav'
gbm_model_v2 = pickle.load(open(file_export_model, 'rb'))
y_pred_gbm_smote = gbm_model_v2.predict(scale_testX)
rpt_metrics_report(y_test, y_pred_gbm_smote, 'GBM - scale and SMOTE ', rpt_confusion_matrix=True)
```



Model: GBM

```
-- Accuracy: 0.8
-- AUC      : 0.7095220490261945
-- Recall   : 0.5127388535031847
```

-- Confusion Matrix

0 FP

FN 1

```
[[1538 159]
 [ 306 322]]
```

-- Metrics report

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.83      | 0.91   | 0.87     | 1697    |
| 1            | 0.67      | 0.51   | 0.58     | 628     |
| accuracy     |           |        | 0.80     | 2325    |
| macro avg    | 0.75      | 0.71   | 0.72     | 2325    |
| weighted avg | 0.79      | 0.80   | 0.79     | 2325    |

Model: GBM - scale and SMOTE

```
-- Accuracy: 0.6619354838709678
-- AUC      : 0.7307992936204393
-- Recall   : 0.8805732484076433
```

-- Confusion Matrix

0 FP

FN 1

```
[[986 711]
 [ 75 553]]
```

-- Metrics report

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.93      | 0.58   | 0.72     | 1697    |
| 1            | 0.44      | 0.88   | 0.58     | 628     |
| accuracy     |           |        | 0.66     | 2325    |
| macro avg    | 0.68      | 0.73   | 0.65     | 2325    |
| weighted avg | 0.80      | 0.66   | 0.68     | 2325    |


### 3rd Insight and feature importance

- The global metric accuracy drop between these 2 models, however RECALL is more important to focus on -> Customer churn
- The recall / True Positive Rate increase from 51,27% to 88,05%

In [19]:

```
# Plot feature importance
feature_importances = pd.DataFrame({'features': features,
                                    'importance': gbm_model_v2.feature_importances_}).sort_values(
    'importance', ascending=False)

ax = sns.barplot(x="importance", y="features", data=feature_importances, palette="coolwarm")
```



## Summary

### Top 3 main important features related to churn (Yes) are:

- Month-to-Month contract
- Small to Medium Tenure and associated with
- Higher Monthly Charges

### Final considerations

**The recall metric was increased in 71,73% between baseline model and the final version - (88,05% versus 51,27%)**

- Using the right technique it is possible to build machine learning models that really help in business. The current model chosen for deployment will provide higher customer base of possible churners to interact with focus on customer retention

In this example of the final GBM model deployed the recall metric is 88,05% and much higher than the first GBM model

**The next and final notebook will resume all information provided by previous notebooks and build the final script for deployment**

In [2]:

```
# !jupyter nbconvert --to html Phase_2_Build_ML_models_with_Python_5x6_evaluation_using_RECALL_metric.ipynb
```

In [ ]: